

UPLB IDSC-FANS KNOWLEDGE HUB: A KNOWLEDGE MANAGEMENT SYSTEM
FOR FOOD AND NUTRITION SECURITY

JOSE GABRIEL F. SANTOS

SUBMITTED TO THE FACULTY OF THE INSTITUTE OF COMPUTER SCIENCE
UNIVERSITY OF THE PHILIPPINES LOS BAÑOS
IN PARTIAL FULFILLMENT OF THE
REQUIREMENTS FOR THE
DEGREE OF

BACHELOR OF SCIENCE
(Computer Science)

MONTH YEAR

APPROVAL SHEET

The thesis hereto attached entitled UPLB IDSC-FANS KNOWLEDGE HUB: A KNOWLEDGE MANAGEMENT SYSTEM FOR FOOD AND NUTRITION SECURITY, prepared and submitted by JOSE GABRIEL F. SANTOS, in partial fulfillment of the requirements for the degree of BACHELOR OF SCIENCE (Computer Science) is hereby accepted.

MEMBER NAME

Member, Guidance Committee

MEMBER NAME

Member, Guidance Committee

CONCEPCION L. KHAN

Chair, Guidance Committee

Accepted in partial fulfillment of the requirements for the degree of BACHELOR OF SCIENCE (Computer Science).

DIRECTOR NAME

Director, Institute of Computer Science

DEAN NAME

Dean, College of Arts and Sciences
University of the Philippines Los Baños

BIOGRAPHICAL SKETCH

JOSE GABRIEL F. SANTOS

ACKNOWLEDGEMENT

This thesis would not have come together without the people who stood by me while I worked on it.

I want to thank my adviser, Associate Professor Concepcion L. Khan, for her guidance and patience throughout this study. She gave me clear feedback when I needed direction and was generous with her time when I had questions, even the small ones. I learned a lot from working with her.

I would also like to thank the Institute of Computer Science at the University of the Philippines Los Baños. The years I spent here gave me the training and the perspective that made this project possible, and I am grateful to the faculty and staff who made that environment what it is.

To the UPLB Computer Science Society (COSS), my home org for most of my time in UPLB, thank you for the friendships and for the community. A lot of the long nights would have been a lot harder without the people I met through COSS, and I am proud to have been part of it.

To my friends, thank you for being there when the work piled up. The conversations, the food runs, and the company kept me sane more times than I can count.

To my family, thank you for your support and for giving me the space to take this on at my own pace. None of this would have been possible without you, and I do not take that for granted.

And to my partner, thank you for sticking with me through the long stretches when I had little to give. Your patience and steady presence carried me through more of this than you probably realize.

To all of you, thank you. This work belongs to you as much as it does to me.

JOSE GABRIEL F. SANTOS

Contents

INTRODUCTION	1
Background of the Study	1
Statement of the Problem	1
Significance of the Study	1
Objectives	2
Scope and Limitations	3
Definition of Terms	3
REVIEW OF LITERATURE	3
METHODOLOGY	5
Technology Stack	6
System Architecture	7
Key System Features	9
Data Collection	11
Testing	12
RESULTS AND DISCUSSION	14
Participant Profile	22
System Usability Scale	23
Task Success and Effort	23
Feature-Specific Satisfaction	24
AI-Assisted Feature Evaluation	25
CONCLUSION	31
RECOMMENDATIONS	32
Recommendations for Features and Fixes to the System	33
Recommendations for the Study	34
APPENDIX A: SURVEY QUESTIONNAIRES	34
Consent and Data Privacy Notice	35
Demographics	36
Screening Item	36
Standard User Track — Task Block	36

AI-Assisted Feature Evaluation Block	38
Public Analytics Task	39
Organization Administrator Track — Task Block	39
Site Administrator Track — Task Block	41
System Usability Scale (SUS)	42
Feature-Specific Satisfaction Grids	42
Open-Ended Items	44

List of Tables

1	Infrastructure	7
2	Elasticsearch Indices	9
3	Analytics Charts (14 Total)	10
4	Evaluation Instruments	20
5	Form Structure (Remote Unmoderated, Role-Branched)	21
6	Realized Participant Demographics ($N = 26$)	26
7	SUS Results (Pooled and Role-Stratified)	27
8	Task Success and SEQ by Task. SEQ is reported on the form's 1 = Very Easy, 5 = Very Difficult scale; lower means indicate easier tasks.	28
9	Feature-Satisfaction Items, Mean (SD) on a 1–5 Likert (5 = Strongly Agree). Each item has $n = 26$; the Admin Tools grid was answered substantively by admin tracks only ($n = 6$).	29
10	AI Insight Likerts (5 = Strongly Agree, $n = 26$ per insight)	29
11	AI Extraction Accuracy per Field. Wilson 95% CI on the N/A-excluded de- nominator; three of 26 participants reported they could not see the extracted metadata after upload and are excluded from the per-field denominator. . . .	30

List of Figures

1	System Architecture	8
2	MongoDB-Elasticsearch Schema Connection	15
3	Home Page	16
4	Organization Hierarchy	17
5	Organization Page	17
6	Papers Page with mark	18
7	Post Page with AI Related Insights	18
8	AI Extraction — Post Submission	19

ABSTRACT

JOSE GABRIEL F. SANTOS, University of the Philippines Los Baños, MONTH YEAR.
**UPLB IDSC-FANS KNOWLEDGE HUB: A KNOWLEDGE MANAGEMENT
SYSTEM FOR FOOD AND NUTRITION SECURITY**

Major Professor: PROF. CONCEPCION L. KHAN

Food and nutrition security research in the Philippines is scattered across different institutions, databases, and studies, with no one unifying system to bring it all together. This makes it harder for researchers, policymakers, and even the general public to find and use the information that is already there. This study presents the UPLB IdSC-FaNS Knowledge Hub, a web-based Knowledge Management System (KMS) built specifically for the UPLB Interdisciplinary Studies Center on Food and Nutrition Security. The system was developed using a hybrid MongoDB and Elasticsearch architecture for fast full-text search and aggregations, a community layer with organizations and a moderated approval pipeline, and AI-assisted curation through Gemini-based metadata extraction and a related-insights panel that synthesizes themes and research gaps from each post. The system was evaluated through a fully remote, unmoderated study on $N = 26$ participants from the UPLB community, split across three tracks: 20 standard users, 5 organization administrators (one of whom was excluded from admin-gated analyses due to a post-hoc-identified provisioning gap, leaving $n = 4$), and 1 website administrator reported as a single-case feasibility observation. The pooled SUS came out to 83.17 (95% CI [76.35, 89.52]), placing the system in the Good–Excellent band. Twenty-two of twenty-three tasks scored $\geq 95\%$ success at SEQ ≤ 1.5 on a 1–5 scale where 1 is easiest, and per-field AI extraction accuracy ranged from 86.4% on Year/Journal to 100% on Title and Abstract. AI metadata extraction was also the most-named useful feature in the open-ended responses. The five research objectives were all met at both the satisfaction and task-success layers, and what is left is mostly a focused bug-fix list (cite-button discoverability, post-submit server validation, sort and filter regressions, and an AI-trust dip addressable through provenance cues) rather than a structural rework.

INTRODUCTION

Background of the Study

Food security and nutrition are essential topics relevant to millions of people especially in a developing country like the Philippines. The Philippine nutrition survey in 2013 shows that over half of the more than twenty thousand participants had inadequate amounts of essential nutrients, with deficiencies in key micronutrients like Iron, Calcium, and Vitamin C affecting 90% of the population. It was also noted that participants in lower socioeconomic groups were more likely to be deficient in the essential nutrients Angeles-Agdeppa et al. (2019). Food insecurity, on another hand, is disproportionately low especially after the COVID-19 pandemic among small scale farmers who were most affected by the disrupted market, larger households who share limited resources with more people, and indigenous people who have limited access to economic and social service Durano et al. (2024).

Research on nutrition is important because it provides an understanding of how nutrition influences overall well-being, aids in disease prevention and management, and addresses the disparity of health among different socioeconomic groups Heitman et al. (2024). These are also some of the reasons why the Interdisciplinary Studies Center on Food and Nutrition Security or IdSC-FaNS was formed. It was created under the University of the Philippines Los Baños Office of the Vice Chancellor for Research and Extension for the continued stewardship of nutrition and food security research in the country Global Nutrition Report (2020).

Statement of the Problem

Data and research on food and nutrition security in the Philippines are distributed across different databases, studies, and institutions with no one unifying system. This can hamper and delay policy makers, other researchers, and the general public in accessing and interpreting information Righettini and Bordin (2022). The lack of access to comprehensive data and analytics for food security and nutrition research can lead to delayed action from policy makers, inadequate public awareness, and limited research collaboration.

This study proposes the development of a Knowledge Management System (KMS) for food and nutrition security research in the Philippines. This aims to create a centralized, accessible platform for information essential to food security research in the Philippines.

Significance of the Study

The Knowledge Management System (KMS) provides the IdSC-FaNS a platform for sharing knowledge on food and nutrition security. Data visualization and analytics on the current

trends in food security research enhance the ability of IdSC-FaNS to monitor, interpret, and address the challenges in food and nutrition security. The use of advanced data visualization and analysis allows the IdSC-FaNS and researchers to identify trends and emerging issues in the field. This supports data-driven decisions, but also encourages targeted research that addresses the current and future issues. The KMS also provides a repository where research documents, reports, reviews, and data can be found. All of the documents are indexed for easy retrieval and access using Elasticsearch. This helps researchers, policymakers, and other stakeholders access important information without navigating multiple sources, ultimately making research on food and nutrition security more accessible and actionable. Beyond storing and searching, the KMS introduces two AI-assisted features that increase its significance. A Gemini-based metadata extractor was added to reduce the effort of researchers in inputting research and to decrease the inaccuracy of the data within the system. A Gemini-generated *related insights* panel synthesizes themes and research gaps given the current topic of the post, showing the unseen connections that keyword search alone cannot reveal. Both features support the thesis’s broader contribution: a domain-focused knowledge management system that integrates full-text search, community discussion, and AI-assisted curation in a unified platform appropriate for a developing-country research context.

Objectives

The main objective of this study is to develop a web application that serves as a Knowledge Management System to provide easy access to food and nutrition security research data. Specifically, it aims to:

1. Provide a platform for research on food and nutrition security where the public can access curated and approved research.
2. Enable advanced search functionality for easy access to research data using Elasticsearch.
3. Provide data visualization and analytics tools that help users identify trends, emerging challenges, and research gaps in food security and nutrition.
4. Support community-driven curation through organizations and moderated discussions.
5. Provide AI-assisted curation through automated metadata extraction and research insight generation.

Scope and Limitations

The study focuses on the design and development of a web-based Knowledge Management System for food and nutrition security in the Philippines. It collects and stores research data; it enables advanced searching and filtering; and it provides advanced data and analytic visualization. The platform relies on the availability of research data from various institutions and organizations affiliated with the ICS-FaNS.

Definition of Terms

- **Food Security** — a state in which all people, at all times, have physical, social, and economic access to sufficient, safe, and nutritious food that meets their dietary needs and food preferences for an active and healthy life.
- **Nutrition Security** — a state where all people have access to essential nutrients and adequate diets that support physical and mental well-being, reduce the risk of disease, and promote a productive life.
- **Knowledge Management System** — a structured system that collects, stores, shares, and analyzes information within an organization.
- **Elasticsearch** — open-source search and analytics engine designed to enable fast, full-text search and data retrieval across large datasets. It is also a datastore and vector database designed for near real-time searches.

REVIEW OF LITERATURE

Knowledge management (KM) is the process of making the right information available to the right people. It is capturing valuable information, making it accessible, and using it to improve the work of organizations Van Zyl et al. (2014). IBM identifies three types of knowledge KM usually works with namely: tacit, implicit, and explicit knowledge. Tacit knowledge comes from experience and thus is difficult to articulate and communicate. Similar to tacit knowledge, implicit knowledge comes from “know-how” knowledge that has not been documented yet. On the other hand, explicit knowledge is structured information easily articulated and documented, and can easily be transferred between people IBM (2023). Knowledge management involves translating tacit knowledge into explicit knowledge and vice versa for easier understanding and an increase in organizational efficiency. In higher education, this can be seen as the transfer of knowledge between research and research teams, ensuring that the findings are effectively disseminated and applied.

A knowledge management system (KMS) is an information system made to facilitate the management of information in an organization. It is a tool and a platform designed to create, collect, organize, store, and disseminate knowledge to the organization. The KMS is then leveraged by the organization in order to improve decision making, and efficiency Maharani et al. (2020). A Knowledge Management System (KMS) typically consists of several key components: databases and repositories for storing knowledge like research papers, case studies, and manuals; search engines that allow users to find specific knowledge quickly; tools that facilitate communication, such as discussion forums and shared documents; data analytics and visualization tools that help users interpret knowledge by providing insights into trends, and gaps in the research field; and knowledge sharing platforms that allow the members to contribute and share their expertise and experiences. These components make sure that knowledge is efficiently organized, accessible, and shareable within the system and the organization Usman et al. (2019).

Several successful knowledge management systems have been developed and implemented. Harvard’s DASH (Digital Access to Scholarship at Harvard) repository serves as an open-access platform where researchers can upload and share their academic work. It uses Elasticsearch for quick and accurate searching across thousands of research papers in their repository Harvard Library (2024). The Food and Agriculture Organization (FAO) has also developed a platform where data and research on food security are kept, and its open-access nature allows governments and policymakers from third world countries to easily view and collaborate with world-class researchers promoting innovation Food and Agriculture Organization (2023).

Despite the great effort researchers and developers have made in leveraging KMS, studies and research on KMS have not been sufficiently explored, especially in third world countries. There is still a lack of understanding on how to efficiently utilize the knowledge in a KMS to increase efficiency in an organization Maharani et al. (2020). However, previous research highlights diversity in focus areas; the development of a particular KMS requires specific needs appropriate to the topic and target user group of the system. Maharani et al. Maharani et al. (2020) suggests the People, Process, Technology (PPT) Framework is crucial to the design and development of the said system. The PPT framework states that the people who use the system need the right skills; the process needs to be properly structured; and the technology should be able to support both. This allows the KMS to function without resistance to change of the people and the process, as well as the technological inefficiencies to significantly affect the system.

Artificial intelligence and machine learning integration can also provide KMS enhanced functionality. Natural Language Processing or NLP can be used to analyze large amounts of text data, and extract key insights or trends related to food security. Other AI and ML tools like AI-powered search engines increase the accuracy and efficiency of the engine Vinutha and Padma (2020). In this paper, a variety of statistical and algorithmic AI and ML techniques were examined for their advantages and disadvantages, most of which proved advantageous to the search engine's relevance scoring and type matching among others. Utilizing the appropriate techniques is consequential to the performance of the search engine and ultimately the KMS being developed.

METHODOLOGY

This chapter outlines the methodology used to design and implement a food and nutrition security research knowledge management system. The primary objective of this project is to design and develop an efficient KMS — UPLB IdSC-FaNS Knowledge Hub, that makes the access and collaboration of knowledge and ideas easier especially for students, researchers, and policymakers.

Technology Stack

The following technologies are used to create the web-based Knowledge Management System:

- **React** — a JavaScript library used for building dynamic user interfaces.
- **Node.js** — a JavaScript runtime environment used for backend API development and other server-side processes.
- **Next.js** — a React framework that enables server-side rendering and static site generation for building web applications.
- **Express.js** — a Node.js web framework for building RESTful APIs.
- **MongoDB** — NoSQL database used for storing structured data.
- **Elasticsearch** — a full-text search engine.
- **TanStack React Query** — an asynchronous state management library used for fetching, caching, and synchronizing server state in React apps.
- **Axios** — a promise-based HTTP client used to make requests to APIs from both the browser and Node.js.
- **Socket.io** — a library that enables real-time, bi-directional, and event-based communication between web clients and servers.
- **Tailwind CSS** — a utility-first CSS framework used for rapidly building custom user interfaces without leaving your HTML.
- **shadcn/ui** — a collection of re-usable, accessible UI components built with Tailwind CSS and Radix UI that you copy and paste into your apps.
- **Recharts** — a composable charting library built on React components, used for rendering data visualizations and graphs.

- **Google Auth** — an authentication service that allows users to sign in to applications using their existing Google account credentials.
- **Gemini 2.5 Flash** — Google’s high-efficiency multimodal AI model, optimized for speed and reasoning across text, images, and code.

Table 1: Infrastructure

Component	Provider	Role
App Platform (sgp1)	DigitalOcean	Frontend + backend hosting
Managed MongoDB	DigitalOcean	Production database
Managed Elasticsearch	DigitalOcean	Production search
Spaces + CDN	DigitalOcean	PDF and image storage
Gemini API	Google	AI extraction and insight generation
CrossRef REST API	CrossRef	DOI-based metadata enrichment

Table 1 shows the infrastructure the KMS is built on. It shows the components, the provider of that component, and what its role is. Most of the components are set up on DigitalOcean. The frontend and the backend are both on separate app platforms.

System Architecture

A few of the main goals of the KMS are to provide easy access to research and knowledge through search and discoverability; the system should allow quick and full-text searching through data. In order to achieve this, the application employs polyglot persistence or the use of multiple databases to support full searching capabilities while keeping storage overhead low and maintaining the integrity of the database.

The KMS uses two different technologies:

- MongoDB stores structured and relational data like Users, Organizations, and others.
- Elasticsearch handles content data, i.e. research papers, posts, and comments for quick searching and retrieval.

MongoDB stores all transactional and relational data: users, organizations, posts, comments, papers, notifications, organization-creation requests, user activities, featured posts, reports, moderation logs, and AI-insight cache. Elasticsearch stores a denormalized projection of posts (`kms_posts`) and papers (`kms_papers`), both of which are optimized for searching. Synchronization from MongoDB to Elasticsearch is handled by Mongoose `post('save')`




Figure 1: System Architecture

and `post('findOneAndUpdate')` hooks defined directly on the `Post` and `Paper` schemas. Mongoose acts as a middleware and checks data synchronization between the two databases.

PDFs and images are stored in **DigitalOcean Spaces** (an S3-compatible object store) and delivered through the Spaces CDN. MongoDB documents persist only the CDN URL, keeping document size small.

A database schema defines how data is structured in the system. Because the KMS features a hybrid architecture with MongoDB and Elasticsearch, the database schema is labelled with two parts: one for MongoDB; one for Elasticsearch. However, both parts still function as one because they are synchronized.

As shown in Figure , the `User`, and `Organization`, are stored and accessed in MongoDB. Posts and papers are indexed in Elasticsearch as these are the searchable data in the application.

MongoDB's `text` index lacks four capabilities required for advanced searching: field-level boosting (e.g., title above body), result highlighting via `<mark>`, fuzzy matching, and

aggregations over full-text matches. Elasticsearch supplies all four via BM25-based relevance scoring on Apache Lucene Robertson and Zaragoza (2009); McCandless et al. (2010). The system therefore maintains two stores and synchronizes them via Mongoose hooks.

The system is a separated frontend and backend. The Next.js frontend is compiled to a production bundle and served from DigitalOcean App Platform. It communicates with the Express backend over a REST API using Axios, with JWT session cookies set as `httpOnly`, `Secure`, and `SameSite=None` to allow cross-origin requests from the authenticated frontend. A persistent Socket.io websocket channel is open to deliver the real-time notifications.

Table 2: Elasticsearch Indices

Index	Source	Denormalized fields	Boosted fields
kms_posts	posts	author name, organization name, tags	title ³ , body ¹ , tags ²
kms_papers	papers	first-author surname, journal	title ³ , abstract ¹ , keywords ² , authors ²

Key System Features

The IdSC FaNS Knowledge Hub is designed to enhance the accessibility, discovery, collaboration, and sharing of knowledge related to food and nutrition security. Below are the key features that make the app an effective platform for researchers, students, and policymakers.

1. **Centralized Repository:** The system allows searching, filtering, and sorting of all published content (posts and papers). Content curation is centered around **Organizations**, which are units of moderation with a four-tier role hierarchy: owner, administrator, member, and follower. Members may submit posts to their organizations; submissions enter a pending state and become published only after an organization administrator approves them, or hidden if rejected with a reason. Once published, papers and posts are viewable to all users regardless of their role in an organization. Depending on a user’s role in the organization the content belongs to, it is also interactable (like, dislike, comment).
2. **Advanced Search:** The global `/search` page presents three tabs (All, Posts, Papers), a left filter panel, paginated results, and marks highlighting matching terms in the results. All searches happen through Elasticsearch with field level per Table 2. The result relevance is BM25 weighted. The home page can also be filtered through eight pre-made topics and categories discussed in **Other AI Enhancements** key feature.

3. **Analytics and Data Visualization:** The admin dashboard keeps track of system-wide information and statistics on users, organizations, papers, and posts. The organization analytics tab, present in every organization page, shows posts over time, member growth, post-type distribution, and top tags of the organization. The site-wide research trends page is open to all registered users and reports research output, research areas, and organization comparisons. All charts are rendered with Recharts from data fetched from Elasticsearch aggregations on the backend.

Table 3: Analytics Charts (14 Total)

Scope	Chart	Data source
Site	Posts over time	<code>kms.posts</code> aggregation
Site	Papers over time	<code>kms.papers</code> aggregation
Site	New users over time	<code>users.createdAt</code> aggregation
Site	Post-type distribution	<code>kms.posts post.type</code> terms agg
Site	Top tags	<code>kms.posts tags</code> terms agg
Site	Top organizations by post count	<code>kms.posts organization</code> terms agg
Org	Posts over time in the org	<code>kms.posts</code> filter + date histogram
Org	Member growth	<code>organizations.members</code> over time
Org	Post-type distribution within org	<code>kms.posts</code> filter + terms agg
Org	Top tags within org	<code>kms.posts</code> filter + terms agg
Public	Research output	<code>kms.papers</code> aggregation
Public	Most active research areas	<code>kms.posts.tags</code> terms agg
Public	Organization comparisons	<code>kms.posts</code> org $\times$ type matrix agg
Per-post	AI-summarized related insights	Gemini + cache

4. **AI Analytical Tools:** AI related insights are rendered on each post page, providing a qualitative analysis of the posts and the research surrounding the post. It provides a Gemini-generated summary of the post, key themes, and research gaps related to the topics of the post. Each generated insight is stored and cached in the database for 7 days to reduce AI consumption.
5. **Other AI Enhancements:** Other AI features are used for the ease of use of content curation, and accuracy of papers uploaded. Each post is automatically categorized into eight different subtopics:
- **Food Security**
 - **Nutrition & Health**

- **Food Science & Technology**
- **Agricultural Systems**
- **Policy & Governance**
- **Traditional Knowledge**
- **Education & Communication**
- **Environment & Sustainability**

This is to assign topics and categorize posts depending on their given topic. Metadata extraction is also done to every research paper being uploaded in the system. This is done through research paper posts by users in their organizations. At the same time, in bulk upload of research papers by organization admins, each paper is extracted for its metadata to be added to the MongoDB and Elasticsearch. Duplicates are screened and skipped in the uploading.

Data Collection

User-Driven PDF Upload

The system is designed for decentralized contribution by IdSC-FaNS members and UPLB constituents rather than bulk ingestion by the researcher. A user uploads a PDF through the Submit a Post dialog, the file is then sent as a request to `POST /api/upload/paper`, and stored in DigitalOcean Spaces. The PDF is encoded and sent to Google Gemini 2.5 Flash with a prompt instructing it to return the metadata of the paper. The extracted metadata prefills the create paper form, which the user reviews, edits, and submits. This is a **human-in-the-loop** design consistent with Human-Centered AI principles Shneiderman (2020).

Admin Pre-seeded Research

For demonstration and evaluation purposes, publicly accessible research provided by the UPLB University Library in the initial data collection phase, and manually chosen papers from the University Knowledge Digital Repository (UKDR), were added to the system through the bulk upload feature as an organization admin. The collection of data involves the gathering of relevant material and making sure the data is structured, formatted, and indexed properly in Elasticsearch.

Research submitted by the testers and evaluators were also added to the system.

Testing

A **convergent mixed-methods evaluation** Creswell and Plano Clark (2017) was adopted, combining quantitative usability and task-success measurements with qualitative thematic feedback collected in the same session. The evaluation operationalizes usability per the **ISO 9241-11:2018** definition as a composite of effectiveness, efficiency, and satisfaction International Organization for Standardization (2018).

Participants and Sampling

The evaluation was administered as a fully remote, unmoderated study Sauro and Lewis (2012); Albert and Tullis (2013). Each participant was given a link to the deployed website and the Google Form and worked through the test at their own pace and on their own machine.

Three roles were created through **role-based form branching**. A screening question (“What role were you assigned for this evaluation?”) routed each participant into one of the three task blocks — Track A (Standard User), Track B-Org (Organization Administrator), and Track B-Site (Website Administrator). The tracks were not randomly assigned; instead, the researcher assigned each participant their task depending on their role in IdSC-FaNS and their relevant qualifications.

The target total was $N \geq 15$, the working floor for stable SUS confidence intervals Lewis (2018). Nielsen’s five-user heuristic Nielsen (2000) and Faulkner’s refinement Faulkner (2003) indicate that 10–15 testers detect $\geq 95\%$ of usability problems, supporting this floor. A total of $N = 26$ responses were collected, distributed as **20 standard users**, **5 organization administrators**, and **1 website administrator**. The total number of valid responses comfortably exceeds the $N \geq 15$ stable-SUS floor, and Track A and Track B-Org each meet or exceed the five-user range at which single-role usability findings stabilize Nielsen (2000); Faulkner (2003). In contrast, Track B-Site did not meet the sample threshold: only one website administrator was available within the evaluation window because the role is provisioned for at most a handful of accounts in the production deployment. Site-admin findings are therefore reported as a **single-case feasibility observation** confirming that the site-administrator workflow is operable end-to-end and are excluded from role-stratified statistical inference; the eligible site-administrator population is itself small (on the order of three accounts in the production deployment), and the inferential follow-up is named in Chapter V.

A second sample-related limitation surfaced after the testing and is recorded here. One Track B-Org participant was assigned the org-administrator role for the evaluation but was

not actually elevated to org-admin permissions in the production system. The participant’s open-ended response described an inability to locate admin controls on the organization page, which the rest of the Track B-Org task data did not align with. Because the system enforces role-based access control, the participant’s failures on admin-gated tasks reflect the absence of the admin role rather than a usability property of the admin tools. The participant is therefore excluded from Track B-Org task and SUS analyses, dropping the role denominator for Track B-Org to $n = 4$ for those analyses; the participant’s responses on items that do not depend on admin privileges (pooled SUS, demographics, AI insight Likerts, AI extraction accuracy, and the public-analytics task) are retained. The reduced denominator is one participant below the five-user floor cited in Nielsen Nielsen (2000) and refined by Faulkner Faulkner (2003). Track B-Org is therefore reported with sample-based descriptive statistics only, while the inferential SUS claim is carried by the pooled $N = 26$ which exceeds the $N \geq 15$ stable-CI floor cited from Lewis Lewis (2018). The case is itself an argument for a synchronous provisioning verification step in any follow-up evaluation, which is discussed in Chapter V.

Evaluation Instruments

SUS remains one of the most widely used tools for measuring usability in HCI because it is short, free, and still dependable even with smaller sample sizes. It also gives a clear 0–100 score that can be interpreted using familiar labels, such as “Good” (around 70–84) and “Excellent” (85 and above) Bangor et al. (2009). Task success was also included to measure effectiveness as stated in ISO 9241-11:2018, International Organization for Standardization (2018). Along with that, Single Ease Question (SEQ) test was also employed to capture perceived task difficulty in a single question making it more practical for the researcher and the evaluation participants alike Sauro and Dumas (2009). Task success and SEQ are done for every task in supplementary with each other to determine the completeness, accuracy, and difficulty each task is. Feature-level Likert ratings were added so the Results chapter can clearly show which areas performed well or need improvement (search, AI, community, analytics, and admin tools) Lewis and Sauro (2018). For AI insights, ratings were completed by domain experts since judging summary quality and research-gap validity requires subject knowledge Shneiderman (2020). Extraction accuracy was evaluated using a per-field information-retrieval approach Hripcsak and Rothschild (2005). Finally, open-ended responses were analyzed using Braun and Clarke’s six-phase thematic analysis to generate evidence-rich insights for the Discussion chapter Braun and Clarke (2006).

The table of all evaluation instruments used is found on the next page. Seven instruments were used. A mix of both qualitative and quantitative methods are optimal to gather complex

but relevant data. Subjective and objective methods were also used to accurately test the features of the system while obtaining anecdotal experiences when using the system.

The test was administered using Google Forms. Evaluators were instructed to have both the form and the website open in tabs side-by-side. They were to strictly follow the given instructions per task before answering the relevant questions. Because the evaluation was self-paced and unmoderated, the protocol is realized as the structure of the Google Form rather than a timed in-person session. The form is a single linear flow: consent and Data Privacy Act notice → demographics (age band, gender, role/affiliation, qualifications, research-platform usage frequency, device used) → screening item (“What role were you assigned for this evaluation?”) → role-branched task block → AI-evaluation block (shared) → public-analytics task (shared) → SUS (10 items) → feature-satisfaction Likert grids (Search and Papers; AI Features; Community and Organization; Analytics; Admin Tools — answered only by admin tracks; Overall) → six open-ended thematic items → submission. Per-task instruments — success outcome (Yes / Partially / No), SEQ difficulty (1–5, where 1 = Very Easy and 5 = Very Difficult), and an optional free-text problem report — are placed immediately below each task instruction so that the rating reflects the just-completed task.

The individual sections of the form are discussed more thoroughly in the Results and Discussion. It is also important to note that for testing, all new accounts were immediately added to the *IdSC-FaNS* organization as members to ease the bottleneck of having to approve the users in real-time.




Figure 2: MongoDB-Elasticsearch Schema Connection




Figure 3: Home Page




Figure 4: Organization Hierarchy




Figure 5: Organization Page




Figure 6: Papers Page with mark




Figure 7: Post Page with AI Related Insights

Organization * Post Type *

Research Paper

Research papers must belong to an organization.

Title *

Body

B *I* H₂

☰

☷

🔗

↶

↷



Upload your research paper PDF

Upload your file and our AI will automatically extract the title, authors, abstract, keywords, and other details for you.




Figure 8: AI Extraction — Post Submission

Table 4: Evaluation Instruments

ID	Instrument	Data type	Scale	Citation
I-1	System Usability Scale	Quantitative	10 items, 1–5 Likert, 0–100	Brooke (1996); Bangor et al. (2009)
I-2	Task success rate	Quantitative	Yes / Partial / No per task	Rubin and Chisnell (2008); International Organization for Standardization (2018)
I-3	Single Ease Question (SEQ)	Quantitative	1–5 per task	Sauro and Dumas (2009); Sauro (2012)
I-4	Feature-specific satisfaction	Quantitative	5-point Likert grids	Lewis and Sauro (2018); Likert (1932)
I-5	AI insight accuracy	Mixed	4×5 -point Likert per insight	Shneiderman (2020); Doshi-Velez and Kim (2017)
I-6	AI extraction accuracy	Quantitative	Per-field Correct / Partial / Incorrect / N/A	Hripcsak and Rothschild (2005)
I-7	Open-ended thematic feedback	Qualitative	6 free-text items, inductive coding	Braun and Clarke (2006)

Table 5: Form Structure (Remote Unmoderated, Role-Branched)

Section	Items	Instruments	Track
Consent & demographics	7 items	—	All
Screening (assigned role)	1 item	—	All
Track A task block	7 task groups (sign-in & profile, profile expertise, search, save/cite/download, follow & join org, like/comment, submit post)	I-2, I-3	Track A
Track B-Org task block	8 task groups (approve, reject, manage join requests, pin, promote, moderate report, bulk upload, org analytics)	I-2, I-3	Track B-Org
Track B-Site task block	5 task groups (approve org request, site analytics, announcement, user management, feature post)	I-2, I-3	Track B-Site
AI-evaluation block	Insight 1 (mandatory) + Insight 2 (optional) + paper upload	I-2, I-3, I-5, I-6	All tracks
Public-analytics task	1 task	I-2, I-3	All tracks
SUS	10 items	I-1	All tracks
Feature-satisfaction Likert grids	22 items across 6 grids (Admin grid answered only by admins)	I-4	All tracks
Open-ended	6 items (most useful, most confusing, missing features, bugs, plan to use again, other comments)	I-7	All tracks

RESULTS AND DISCUSSION

This chapter reports the results of the evaluation defined in Chapter III. It interprets the data collected against the five research objectives stated in Chapter I. The sample data comprises $N = 26$ participants: 20 standard users, 5 organization administrators, and 1 website (super) administrator. This chapter reports the participant’s profile, the data and results from the System Usability Scale, per-task success and effort, feature-specific satisfaction, AI-assisted feature evaluation, and the qualitative evaluation in that order.

Participant Profile

The 26 participants spanned four age bands and two affiliations of interest. Twelve (46%) were 18–24, nine (35%) were 25–34, four (15%) were 45–54, and one (4%) was 55+. Eighteen (69%) identified as female and eight (31%) as male. Reported affiliation was distributed across “Other” (15; 58%), Undergraduate Student (4; 15%), Faculty (3; 12%), Research Staff (3; 12%), and Administrative Staff (1; 4%). Thirteen participants listed an additional research qualification as Registered Nutritionist-Dietitian and a further five listed graduate or professional credentials (PhD, MSc $\times 2$, Registered Agriculturist, DMD); the remaining seven reported none. Research platform usage frequency of the users are around the higher end of 1–5 scale. Ten participants report 3, eight reported 4, four report 5, and four reported 2; no participant reported 1. Twenty-five participants (96%) tested on a laptop or desktop and one used a smartphone.

All but one of the participants who answered “Other” in their UPLB role are UPLB Alumni who graduated with a Degree in BS Nutrition under the Institute of Human Nutrition and Food.

Crossing affiliation against assigned evaluation role shows that the five organization administrators were drawn from “Other” (4) and Undergraduate Student (1), and the single website administrator was drawn from “Other”; the standard-user track included all three Faculty respondents and all three Research Staff respondents. This distribution is consistent with the role-provisioning approach described in Chapter III: admin accounts were elevated by the researcher before recruitment and so the admin tracks reflect *who was given an admin account*, not *who happened to be senior in real life*. As noted in Chapter III, one Track B-Org participant was assigned the org-administrator role but was not elevated in the production system, and is excluded from admin-gated analyses for the reasons given there. The Track B-Org admin-track denominator is therefore $n = 4$ for SUS and per-task analyses; the participant’s responses are retained in the pooled SUS and in analyses not contingent on admin provisioning (AI insight Likerts, AI extraction accuracy, public-analytics task).

System Usability Scale

The pooled SUS score across all 26 participants was **mean** = 83.17, **SD** = 17.61, **median** = 92.5, **range** 50.0–100.0. The 95% bootstrap CI (10,000 resamples) was [76.35, 89.52]. The mean falls in the “**Excellent**” band on the Bangor adjective scale (≥ 85 in Bangor’s strict cut, “Good” between 72 and 85) Bangor et al. (2009). Even the lower bound of the bootstrap CI sits inside the “Good” band. Sixteen of twenty-six participants (62%) scored in “Excellent” (≥ 85), three (12%) in “Good” (72–84), five (19%) in “OK” (52–71), and two (8%) in “Poor” (38–51). No participant scored below 50.

By track, **Track A** ($n = 20$) averaged 82.75 (SD = 18.44, median 92.5, range 50–100); **Track B-Org** ($n = 4$, **after excluding the unprovisioned participant per Chapter III**) averaged 89.4 (SD ≈ 6.6 , median ≈ 91.25 , range ≈ 80 –95); and **Track B-Site** ($n = 1$) scored 95.0. After the provisioning-related exclusion, the role-stratified Track B-Org SUS rises into the Good–Excellent band, which is consistent with Track A’s distribution and supports the claim that the admin-side workflows do not depress perceived usability when admin permissions are correctly provisioned. P21’s SUS score is retained in the pooled grid because SUS does not depend on admin permissions; only the Track B-Org role-stratified line excludes it. The website administrator’s single score is reported descriptively per the limitation in Chapter III as a single-case feasibility observation and is not used for statistical inference.

Task Success and Effort

Per-task success was coded $(Yes + 0.5 \cdot Partial) / total \cdot 100$ Sauro (2011). Effort was reported on the form’s 1–5 SEQ scale where 1 = **Very Easy** and 5 = **Very Difficult**, so lower means easier. Tasks restricted to a track use that track’s n as the denominator; the AI-evaluation, paper-upload, and public-analytics tasks were given to all tracks and use $n = 26$. Track B-Org tasks use $n = 4$ following the exclusion of the unprovisioned participant identified in Chapter III; the original responses for that participant on admin-gated tasks reflect the absence of the org-admin role rather than a usability property of the admin tools and are therefore not included in the per-task denominators below.

Twenty-two of the twenty-three tasks scored 95% or higher on success and at or below 1.5 on SEQ. Two tasks remain below the threshold after the Track B-Org exclusion:

- **A4 (Save paper / cite / download)** — 87.5% success (15 Yes, 5 Partial), mean SEQ 1.25. The qualitative reports show this is from the cite-button discoverability problem, not an actual failure of save or download.

- **A7 (Submit a discussion post)** — 80.0% success (16 Yes, 0 Partial, 4 No), mean SEQ 1.40. Reports describe an “internal server error” on submit. One participant found that adding a text body to the post body field let the submission go through, which points to a server-side validation rule that the client does not enforce.

After excluding the unprovisioned Track B-Org participant, every B-Org task with a sample-based denominator returns to or near 100% success. B-Org tasks 3, 4, 5, 6, and 8 (manage join requests; pin; promote; moderate report; view org analytics) each scored 100% (4/4), with mean SEQ values that drop below the 1.5 threshold once the unprovisioned participant’s high-effort responses are excluded. B-Org 7 (bulk upload) is reported with the corrected denominator in Table 8; the residual partial responses on that task are attributable to “some files failed” and “duplicate” notices, both of which are correct system behaviour rather than usability faults. With this correction, no task with a sample-based denominator crosses the 1.5 SEQ threshold or the 3.0 difficulty threshold flagged in the data-analysis plan.

Branched yes/no items confirm the same picture. Of the eighteen Track A participants who reached the post-submission confirmation, eighteen rated “Yes, very clear” that the post had been submitted *for review* rather than published immediately, and two answered “I didn’t notice”. This shows that Objective 5’s approval pipeline is communicated correctly when the submit actually succeeds. Sixteen of twenty Track A search responses described the search results as “Yes, very relevant”, three as “Somewhat relevant”, and one as “Not relevant”.

Feature-Specific Satisfaction

Five Likert grids and one Overall grid were given (1–5 Likert, 5 = Strongly Agree). Each item had $n = 26$. The Admin Tools grid was answered by all participants, but the instructions only asked the admin tracks to give a substantive rating, so its means must be read with that in mind. The $n = 26$ distribution is bimodal (12 “Strongly Agree”, 11 “Neutral”, 3 “Agree” on item 121), which is what happens when non-admins default to “Neutral”.

A few things are worth pointing out. Every grid except AI Features and Admin Tools sits at or above 4.46. The **AI-trust item** (3.81 ± 0.94) is the only sub-4 item in the entire instrument, which lines up with the qualitative theme that AI insights are valued but not blindly trusted. The Overall grid is the strongest of the six: “valuable for the UPLB research community” reaches 4.92 with the lowest dispersion ($SD = 0.27$), which supports the system-level claim even with the AI-trust dip.

AI-Assisted Feature Evaluation

Two AI insights were rated by every participant (the second insight had been an “if time” item but the data shows all 26 responded). Each insight was rated on four dimensions: summary accuracy, theme correctness, gap validity, and usefulness. Means landed between 3.88 and 4.12 with fairly high SDs (1.21–1.31). The high SDs come from a small minority of participants who rated low.

The spread here is what’s worth looking at. A subset of participants rated the AI insights “1” on summary accuracy, which lines up with two of the qualitative responses. One participant wrote that “some AI summary [is] not mentioned in the [post]” and guessed the model was “hallucinating”, and another wrote “Just not reliant on AI doing my research”. Since the AI-insight evaluation is a domain judgment about food and nutrition security research Shneiderman (2020), this kind of spread is expected. Domain experts disagree about the model’s gap claims, and the system’s response to that disagreement (the “Was this useful?” feedback control on each insight) becomes a feature that can be evaluated in any future iteration.

For paper-metadata extraction, every field used a per-field correctness category with an explicit “N/A” option to mark “I couldn’t see the metadata”. Three of the 26 participants reported they could not see the extracted metadata after upload (Track A: row 10 “File was successfully uploaded but was not able to see extracted metadata”). These responses are shown as “N/A” in the per-field counts and are excluded from the percentage denominators below. They are reported separately as a $3/26 = 11.5\%$ upload-rendering failure rate, which fed the 88.5% task-success rate for Shared-2.

The single Authors miss was rated “Partially Correct”. On Year / Journal, one paper had no journal field in the source PDF (correctly flagged as “Not present in the paper”), one was rated “All incorrect”, and one was rated “Some correct”. The aggregate “Overall, the AI metadata extraction was accurate enough to be useful without major manual corrections” Likert came back at **mean** = 4.54, **SD** = 0.81 ($n = 26$). This matches the per-field numbers and supports the human-in-the-loop framing of the extraction pipeline in Chapter III Shneiderman (2020).

Table 6: Realized Participant Demographics ($N = 26$)

Variable	Level	Count	%
Age band	18–24	12	46
	25–34	9	35
	45–54	4	15
	55+	1	4
Gender (assigned at birth)	Female	18	69
	Male	8	31
UPLB role / affiliation	Other	15	58
	Undergraduate Student	4	15
	Faculty	3	12
	Research Staff	3	12
	Administrative Staff	1	4
Additional qualifications (free-text)	RND	13	50
	PhD	2	8
	MS / MSc	2	8
	R.Agr. / DMD	2	8
	None / N/A	7	27
Research-platform usage (1 rare – 5 daily)	3	10	38
	4	8	31
	5	4	15
	2	4	15
Device	Laptop / Desktop	25	96
	Smartphone	1	4
Assigned evaluation role	Standard User	20	77
	Organization Admin	5	19
	Website / Super Admin	1	4

Table 7: SUS Results (Pooled and Role-Stratified)

Stratum	n	Mean	SD	Median	Range	95% Bootstrap CI	Adjective band
Pooled	26	83.17	17.61	92.5	50.0–100.0	[76.35, 89.52]	Good–Excellent
Track A (Standard User)	20	82.75	18.44	92.5	50.0–100.0	—	Good–Excellent
Track B-Org (Org Admin)	4	89.4	6.6	91.25	80.0–95.0	—	Good–Excellent
Track B-Site (Site Admin)	1	95.00	—	—	—	—	Excellent ($n = 1$, single-case)

Table 8: Task Success and SEQ by Task. SEQ is reported on the form’s 1 = Very Easy, 5 = Very Difficult scale; lower means indicate easier tasks.

Task ID	Description	Track(s)	n	Success %	Y / P / N	Mean SEQ
A1	Sign-in & report pa- pers count	A	20	100.0	20 / 0 / 0	1.20
A2	Update profile exper- tise	A	20	100.0	20 / 0 / 0	1.00
A3	Search for a paper	A	20	100.0	20 / 0 / 0	1.05
A4	Save / cite / down- load	A	20	87.5	15 / 5 / 0	1.25
A5	Follow & join organi- zation	A	20	100.0	20 / 0 / 0	1.05
A6	Like or comment	A	20	97.5	19 / 1 / 0	1.10
A7	Submit a discussion post	A	20	80.0	16 / 0 / 4	1.40
Shared-1	Read AI insight (In- sight 1)	All	26	100.0	26 / 0 / 0	—
Shared-2	Upload a PDF and review extracted metadata	All	26	88.5	23 / 0 / 3	1.27 (overall)
Shared-3	View public analytics	All	26	100.0	26 / 0 / 0	1.27 (overall)
B-Org 1	Approve a pending post	B-Org	4	100.0	4 / 0 / 0	1.00
B-Org 2	Reject a pending post	B-Org	4	100.0	4 / 0 / 0	1.20
B-Org 3	Manage a join re- quest	B-Org	4	100.0	4 / 0 / 0	1.25
B-Org 4	Pin a post	B-Org	4	100.0	4 / 0 / 0	1.50
B-Org 5	Promote a member	B-Org	4	100.0	4 / 0 / 0	1.50
B-Org 6	Moderate a reported post	B-Org	4	100.0	4 / 0 / 0	1.50
B-Org 7	Bulk upload a folder	B-Org	4	87.5	3 Y, 1 P (some failed / duplicates)	1.25
B-Org 8	View organization analytics	B-Org	4	100.0	4 / 0 / 0	1.25
B-Site 1	Approve organization-creation request	B-Site	1	100.0	1 / 0 / 0	1.00
B-Site 2	Site analytics	B-Site	1	100.0	1 / 0 / 0	1.00
B-Site 3	Site-wide announce- ment	B-Site	1	100.0	1 / 0 / 0	2.00
B-Site 4	User management	B-Site	1	100.0	1 / 0 / 0	1.00
B-Site 5	Feature a post	B-Site	1	100.0	1 / 0 / 0	1.00

Table 9: Feature-Satisfaction Items, Mean (SD) on a 1–5 Likert (5 = Strongly Agree). Each item has $n = 26$; the Admin Tools grid was answered substantively by admin tracks only ($n = 6$).

Cluster	Item	Mean	SD
Search and Papers	Search helped me find relevant research	4.88	0.33
	Filters and sorting were useful	4.85	0.46
	Citation generator is useful	4.69	0.68
	PDF download was fast and straightforward	4.96	0.20
AI Features	Metadata extraction would save time vs. manual entry	4.77	0.59
	AI research insights are accurate and valuable	4.62	0.64
	Research gaps identified are worth investigating	4.50	0.65
	I would trust AI features in my own research workflow	3.81	0.94
Community and Organization	Post and comment system supports academic discussion	4.88	0.33
	Organization structure is clear and well-organized	4.69	0.55
	Approval workflow ensures quality without burden	4.69	0.62
	Notifications kept me informed	4.46	0.81
Analytics	Charts help understand current state of research	4.88	0.43
	Analytics useful for identifying research priorities	4.88	0.43
Admin Tools (admins, $n = 6$)	Org admin tools are easy to use	4.04	0.96
	Site admin dashboard provides useful oversight	4.12	0.99
	Bulk paper upload is practical for building libraries	4.00	0.94
	Role-based permission model feels appropriate	4.08	0.98
Overall	I can see myself / my org using this regularly	4.77	0.51
	This system would be valuable for the UPLB research community	4.92	0.27
	The website feels reliable and responsive	4.88	0.33
	The visual design is clean and professional	4.92	0.27

Table 10: AI Insight Likerts (5 = Strongly Agree, $n = 26$ per insight)

Dimension	Insight 1 mean (SD)	Insight 2 mean (SD)
Summary accuracy	3.88 (1.28)	4.00 (1.26)
Theme correctness	4.00 (1.30)	4.04 (1.28)
Gap validity	4.12 (1.21)	4.04 (1.25)
Usefulness	4.12 (1.31)	4.00 (1.26)

Table 11: AI Extraction Accuracy per Field. Wilson 95% CI on the N/A-excluded denominator; three of 26 participants reported they could not see the extracted metadata after upload and are excluded from the per-field denominator.

Field	Correct / n	% Correct	Wilson 95% CI
Title	23 / 23	100.0	[85.7, 100.0]
Authors	22 / 23	95.7	[79.0, 99.2]
Abstract	23 / 23	100.0	[85.7, 100.0]
Keywords (all relevant)	21 / 23	91.3	[73.2, 97.6]
Year / Journal (all correct)	19 / 22	86.4	[66.7, 95.3]

CONCLUSION

This study set out to design and develop UPLB IdSC-FaNS Knowledge Hub: A Knowledge Management System (KMS) for food and nutrition security research in the Philippines, built specifically for the UPLB Interdisciplinary Studies Center on Food and Nutrition Security (IdSC-FaNS). The system was developed as a web application using a hybrid MongoDB and Elasticsearch architecture, with AI-assisted features for paper metadata extraction and research insights, and a community layer for posts, comments, and organization and site level moderation. The evaluation in Chapter IV used a fully remote, unmoderated form on $N = 26$ participants drawn from the UPLB community across three tracks: 20 standard users, 5 organization administrators, and 1 website administrator.

The pooled SUS score of 83.17 (95% CI [76.35, 89.52], median 92.5) places the system in the Good–Excellent band of the Bangor adjective scale. Twenty of twenty-three tasks scored $\geq 95\%$ success and ≤ 1.5 SEQ on a 1 = Very Easy / 5 = Very Difficult scale. AI metadata extraction was the most-named useful feature in the open-ended responses, and field-level extraction accuracy ranged from 86.4% on Year / Journal to 100% on Title and Abstract. The five research objectives stated in Chapter I are revisited below against the data from Chapter IV.

Objective 1 — Provide a centralized platform for IdSC-FaNS research. This objective is met. The four Search and Papers Likerts landed between 4.69 and 4.96, the public-analytics task succeeded for all 26 participants, and the qualitative responses (theme T7) collected unprompted system level endorsements naming the hub as something participants wished had existed during their own research. The main bug that challenged this objective is the cite button disappearance. A straightforward fix is expected and does not threaten the structure and integrity of the system.

Objective 2 — Enable advanced search and discovery using Elasticsearch. This objective is met without issue. Every Track A search task scored 100% success with SEQ ≤ 1.05 , and sixteen of twenty participants rated the search results as “very relevant”. The two filter regressions in theme T5 (oldest-first ordering and the year-range filter) are scoped, in-scope fixes and do not put the BM25-and-aggregation approach into question.

Objective 3 — Provide data visualization and analytics tools. This objective is met without issue. Both Analytics Likerts hit 4.88, the public-analytics task succeeded for all 26 participants, and no participant raised concerns about the chart set in the open-ended items.

Objective 4 — Support community-driven curation through organizations and moderated discussions. This objective is met. The four Community and Organization

Likerts landed between 4.46 and 4.88, the like-or-comment task succeeded for 19 of 20 participants, and all eighteen Track A participants who reached the post-submission confirmation rated “Yes, very clear” that the post had been submitted *for review* rather than published immediately. After the provisioning-related Track B-Org exclusion described in Chapter III, every admin-gated B-Org task (3 through 6 and 8) returns to 100% success on the corrected $n = 4$ denominator, indicating that the admin-side moderation workflows are usable when admin permissions are correctly provisioned. The post-submit server error (theme T4) is a scoped fix and does not put the organization-and-moderation curation model into question.

Objective 5 — Provide AI-assisted curation through automated metadata extraction and research insight generation. This objective is met. AI metadata extraction was the most-named useful feature in the open-ended responses, per-field extraction accuracy ranged from 86.4% on Year / Journal to 100% on Title and Abstract, and the aggregate “accurate enough to be useful without major manual corrections” item came back at 4.54 ± 0.81 . The AI-trust item (3.81 ± 0.94) is the only sub-4 item in the entire feature-satisfaction instrument and lines up with theme T2’s “hallucination” concern; the fix is provenance and verification cues rather than a pipeline redesign.

Across all five objectives, the system meets each one at both the satisfaction and task-success layers. What remains is a bug-fix list (cite-button affordance, post-body server validation, and filter regressions), not a structural rework. The most consistent qualitative finding across affiliations and tracks is that participants would use the system again: twelve out of twenty-two substantive responses planned to explore further and zero declined outright.

The UPLB IdSC-FaNS Knowledge Hub meets the five research objectives stated in Chapter I at both the satisfaction and task-success layers, with a focused bug-fix list rather than a structural rework. The AI metadata extraction was the most-named useful feature in the open-ended responses, the search and analytics features cluster near the ceiling on the Likert grids, and the community layer was endorsed as a tool participants would use again. The strongest single qualitative claim, repeated across affiliations and tracks is that the system would be valuable for the UPLB research community, with several participants stating they wished the platform had existed during their own research. The recommendations listed above are the next concrete steps for taking the system from a thesis prototype to a production tool that IdSC-FaNS can hand to its researchers, students, and policymakers.

RECOMMENDATIONS

The recommendations below are split into two groups. The first set covers features and fixes to the system itself, drawn from the bugs and feature requests named directly by participants in the open-ended items. The second set covers recommendations for future studies on the system.

Recommendations for Features and Fixes to the System

1. **Restore the cite-button affordance on the paper view.** Five participants reported the cite button was missing or hard to find on Task A4, and the 87.5% success rate on that task is fully explained by this. The control should be surfaced on the paper view next to the existing copy-link, save, and download controls.
2. **Fix the post-submit “internal server error”.** Three to four participants hit an internal server error on Task A7. P5 identified the workaround as adding a text body to the post body field, which means the server is enforcing a non-empty body field that the client does not enforce. The fix is to either enforce the body field on the client or relax the server validation, depending on which behavior is intended.
3. **Fix the oldest-first / newest-first sort and the year-range filter.** P17 reported that the “Oldest first” sort did not return the oldest paper first, and P21 reported that a 2025–2025 year-range query returned posts from 2026. These reports concern the public papers page and are independent of the provisioning issue noted in Chapter III; they are likely syntax inaccuracies in the backend query construction and are an easy fix.
4. **Investigate the metadata-not-rendered-after-upload case.** Three of 26 participants reported that the upload succeeded but the extracted metadata was not visible. This is the 11.5% upload-rendering failure rate that drove the 88.5% task-success rate on Shared-2. The cause should be checked at the upload pipeline and the post-upload paper-view render.
5. **Add insight proofs and verification cues.** The AI-trust item (3.81 ± 0.94) was the only sub-4 item in the entire feature-satisfaction instrument, and the qualitative responses (theme T2) named “hallucination” as the concern. Surfacing the source posts or papers behind each insight, and showing which parts of the insight came from which source, would address this without changing the underlying extraction or insight

pipelines. Altering the Gemini prompt to a more appropriate and more encompassing prompt may fix this issue.

Recommendations for the Study

1. **Re-test the org-join-approval and admin-elevation flows end-to-end.** Auto-membership in IdSC-FaNS and pre-elevated admin privileges short-circuited these two real-world workflows during the evaluation. A future evaluation should test them end-to-end from a freshly registered standard-user account.
2. **Run a follow-up evaluation with a larger admin sample.** Track B-Org rests on $n = 5$ and Track B-Site is a single observation. A follow-up evaluation with more organization administrators and at least one additional website administrator would tighten the role-stratified SUS estimates and let the admin-tools cluster (currently 4.00–4.12) be evaluated on a sample-based denominator rather than a mixed $n = 26$ grid.
3. **Expand the data ingestion pipeline beyond the initial UPLB Library dataset.** The current corpus relies on the metadata and DOI-based enrichment of the initial UPLB Library export. Connecting the ingestion pipeline to additional sources (other UP units, partner institutions, and open-access repositories like FAO) would strengthen the centralized-platform claim over time.

Overall, a moderated evaluation session may be more appropriate to more extensively test all workflows. The current setup of a remote, unmoderated evaluation showed weaknesses in its capability to evaluate admin approval workflows: a moderator would also have caught the missing org-admin elevation on the affected Track B-Org account at the start of the session, before any task data was collected, rather than after the fact through the open-ended responses. A synchronous provisioning-verification step is therefore an explicit requirement for any follow-up evaluation. Expanding the realized samples for Track B-Org and especially Track B-Site would also be beneficial in testing the admin controls present in the system. It is important to note that none of the recommendations require a system architecture and workflow overhaul. Simple polishing and streamlining would be enough to address the qualitative comments and feedback and should reflect positively in the quantitative assessments.

APPENDIX A: SURVEY QUESTIONNAIRES

The evaluation instrument was administered as a single Google Form titled *UPLB IDSC-FaNS Knowledge Hub — Evaluation & Feedback*. The form is reproduced below verbatim, grouped by section in the order presented to participants. The form is role-branched: Standard User Track items were shown only to participants assigned a standard-user role, while Organization Administrator Track and Site Administrator Track items were shown only to participants assigned the respective administrator roles. The AI-Assisted Feature Evaluation block, the Public Analytics task, the System Usability Scale, the feature-specific satisfaction grids, and the open-ended items were shown to all participants regardless of track.

Consent and Data Privacy Notice

This study is conducted in accordance with the Data Privacy Act of 2012 (Republic Act No. 10173). The data controller is the researcher, Jose Gabriel F. Santos (jfsantos3@up.edu.ph), under the academic supervision of thesis adviser Conception L. Khan of the UPLB Institute of Computer Science.

I voluntarily agree to participate in the evaluation of the UPLB IDSC-FaNS Knowledge Hub for a CS190 thesis project at the University of the Philippines Los Baños. I understand that my responses will be used only for academic purposes, including thesis documentation, presentations, and possible publication. Results will be reported in summary form and will not identify me personally.

My email address will be collected only to send the response-edit link and for possible follow-up, and it will not be published. Only the researcher and authorized academic reviewers (such as the thesis adviser and panel) may access raw responses, which will be stored securely in password-protected accounts and devices. Raw data will be retained only for the duration of the thesis review process and up to one (1) year after acceptance, after which it will be securely deleted.

I may skip any question or stop participating at any time without penalty. If I decide to withdraw my response after submission, I may contact the researcher to request removal within the study data-cleaning period. No personal or sensitive information beyond my role and field of expertise is being requested, and I may contact the researcher at any time for questions about the study.

Under RA 10173, I also retain the right to access, correct, or object to the processing of my data, and to file a complaint with the National Privacy Commission if I believe my rights have been violated.

By checking the box below, I confirm that I am at least 18 years of age and freely give

my consent to participate under the terms above.

Demographics

1. What is your age range?
2. Gender assigned at birth.
3. What is your role or affiliation in UPLB?
4. What are your other qualifications? (e.g. PhD, MD, MSc, RND). Write N/A if not applicable.
5. How often do you use online research platforms (e.g., Google Scholar, Scopus, ResearchGate, ScienceDirect, institutional repositories)?
6. What device are you using when accessing the website for this evaluation?

Screening Item

1. What role were you assigned for this evaluation? If none was assigned to you, select Standard User.

Standard User Track — Task Block

After performing each assigned task on the system, evaluators were asked the following per-task instruments. Task-specific verification questions, where present, are listed inline. Success was rated on a three-point scale (Yes / Partially / No); difficulty was rated using the Single Ease Question (SEQ) on a 1–5 scale where 1 = Very Easy and 5 = Very Difficult.

Task 1.

1. Were you able to complete this task?
2. How difficult was this task?
3. How many research papers are currently in the system?
4. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 2.

1. Were you able to complete this task?

2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 3.

1. Were you able to complete this task?
2. How difficult was this task?
3. Were the search results relevant?
4. What search term did you use?
5. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 4.

1. Were you able to complete the tasks?
2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 5.

1. Were you able to complete the tasks?
2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 6.

1. Were you able to complete the tasks?
2. How difficult or confusing were the task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 7.

1. Were you able to submit the post?
2. How difficult or confusing was the task?
3. Was it clear that your post was submitted for review, not immediately published?
4. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

AI-Assisted Feature Evaluation Block

This block was shared across all tracks. Likert items were rated on a 5-point scale (Strongly Disagree to Strongly Agree).

AI Insights — Reading.

1. Were you able to find and read the AI insights?

AI Insight 1 (Mandatory).

1. The AI summary accurately describes the main points of the post.
2. The key themes identified by the AI are appropriate and correct.
3. The research gaps identified by the AI are valid and relevant to food and nutrition security research.
4. The AI insight is useful for understanding the broader research context.
5. Anything the AI got wrong, missed, or misrepresented? (OPTIONAL)

AI Insight 2 (Optional).

1. The AI summary accurately describes the main points of the post.
2. The key themes identified by the AI are appropriate and correct.
3. The research gaps identified by the AI are valid and relevant to food and nutrition security research.
4. The research gaps identified by the AI are valid and relevant to food and nutrition security research.
5. Anything the AI got wrong, missed, or misrepresented? (OPTIONAL)

AI Metadata Extraction — Paper Upload.

1. Were you able to upload a paper and see the extracted metadata?
2. Was the extracted TITLE correct?
3. Were the extracted AUTHORS correct?
4. Was the extracted ABSTRACT correct?
5. Were the extracted KEYWORDS correct?
6. Were the extracted YEAR / JOURNAL correct?
7. Overall, the AI metadata extraction was accurate enough to be useful without major manual corrections.
8. Any specific extraction errors worth noting? (OPTIONAL)

Public Analytics Task

1. Were you able to complete this task?
2. How difficult or confusing was the task?
3. What was the most common research topic shown on Analytics?
4. The analytics visualization help identify trends and gaps in the UPLB IDSC-FaNS Knowledge Hub.

Organization Administrator Track — Task Block

After performing each assigned task on the system, evaluators were asked the following per-task instruments. Task-specific verification questions, where present, are listed inline.

Task 1.

1. Were you able to complete this task?
2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 2.

1. Were you able to complete this task?

2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 3.

1. Were you able to complete this task?
2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 4.

1. Were you able to complete this task?
2. How difficult was this task?
3. If you rejected a post, provide a short reason.
4. Was it clear how to find and manage pending posts?

Task 5.

1. Were you able to complete this task?
2. How difficult was this task?
3. Did you encounter any problems or difficulties when completing this task? (OPTIONAL)

Task 6.

1. Were you able to complete this task?
2. How difficult was this task?
3. Did you hide the post, delete the post, or dismiss the report?

Task 7 — Bulk Paper Upload.

1. Were you able to complete the bulk upload?
2. How difficult was this task?

3. The bulk upload process was straightforward.
4. The AI correctly extracted metadata for the uploaded papers.
5. Were any papers flagged as duplicates, and was the dialog clear?

Task 8.

1. Were you able to complete this task?
2. How difficult was this task?

Site Administrator Track — Task Block

After performing each assigned task on the system, evaluators were asked the following per-task instruments. Task-specific verification questions, where present, are listed inline.

Task 1.

1. Were you able to complete this task?
2. How difficult was this task?
3. Which organization was the most active?

Task 2.

1. Were you able to complete this task?
2. How difficult was this task?
3. The announcement feature feels adequate for platform engagement.

Task 3.

1. Were you able to complete this task?
2. How difficult was this task?
3. The user management tools are adequate in managing users.

Task 4.

1. Were you able to complete this task?
2. How difficult was this task?

Task 5.

1. Were you able to complete this task?
2. How difficult was this task?

System Usability Scale (SUS)

All items rated on a 5-point Likert scale: 1 = Strongly Disagree, 2 = Disagree, 3 = Neutral, 4 = Agree, 5 = Strongly Agree.

1. I think that I would like to use this system.
2. I found the system unnecessarily complex.
3. I thought the system was easy to use.
4. I think that I would need the support of a technical person to be able to use this system.
5. I found the various functions in the system were well integrated.
6. I thought there was too much inconsistency in this system.
7. I would imagine that most people would learn to use this system very quickly.
8. I found the system very cumbersome to use.
9. I felt very confident using the system.
10. I needed to learn a lot of things before I could get going with this system.

Feature-Specific Satisfaction Grids

All items rated on a 5-point Likert scale (Strongly Disagree to Strongly Agree).

Search and Papers.

1. The search feature helped me find relevant research easily.
2. The filters and sorting options were useful for narrowing down results.
3. The citation generator is useful for research.
4. Downloading paper PDFs was fast and straightforward.

AI Features.

1. The AI powered paper metadata extraction would save significant time compared to manual entry.
2. The AI research insights provide accurate and valuable context about research topics.

3. The research gaps identified by the AI are worth investigating.
4. I would trust the AI features enough to use them in my own research workflow.

Community and Organization.

1. The post and comment system supports meaningful academic discussion.
2. The organization structure is clear and well-organized.
3. The post approval workflow ensures content quality without being over burdensome.
4. Notifications helped me stay informed about relevant activity.

Analytics.

1. The analytics charts helped me understand the current state of research in the system.
2. The analytics would be useful for identifying research priorities.

Admin Tools (answered only by participants assigned an admin role).

1. The organization admin tools are easy to use.
2. The website admin dashboard provides useful oversight of the platform.
3. The bulk paper upload feature would be practical for building an organization's paper library.
4. The role-based permission model feels appropriate for a research platform.

Overall.

1. I can see myself or my organization using this system regularly.
2. This system would be valuable for the UPLB research community.
3. The website feels reliable and responsive.
4. The visual design is clean and professional.

Open-Ended Items

1. What feature did you find most useful or impressive?
2. What was the most confusing or frustrating part of using the system?
3. Are there features or tools you expected the system to have but are missing?
4. Did you encounter any bugs or unexpected behavior? Please describe.
5. Do you expect to use this system again on your own time?
6. Any other comments, suggestions, or concerns?

LITERATURE CITED

- Albert, W. and Tullis, T. (2013). *Measuring the User Experience: Collecting, Analyzing, and Presenting Usability Metrics*. Morgan Kaufmann, 2nd edition.
- Angeles-Agdeppa, I. et al. (2019). Food sources, energy and nutrient intakes of adults: 2013 Philippines National Nutrition Survey. *Nutrition Journal*, 18(1).
- Bangor, A., Kortum, P., and Miller, J. (2009). Determining what individual SUS scores mean: Adding an adjective rating scale. *Journal of Usability Studies*, 4(3):114–123.
- Braun, V. and Clarke, V. (2006). Using thematic analysis in psychology. *Qualitative Research in Psychology*, 3(2):77–101.
- Brooke, J. (1996). SUS: A quick and dirty usability scale. In Jordan, P. W., Thomas, B., Weerdmeester, B. A., and McClelland, A. L., editors, *Usability Evaluation in Industry*, pages 189–194. Taylor and Francis.
- Creswell, J. W. and Plano Clark, V. L. (2017). *Designing and Conducting Mixed Methods Research*. SAGE Publications, 3rd edition.
- Doshi-Velez, F. and Kim, B. (2017). Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608*.
- Durano, F. L., Sarmiento, J. M., Dee, G., Wahing, G. D., Traje, A. M., and Baker, D. (2024). Food security of agri-food system actors during the COVID-19 pandemic in the Philippines: Post-pandemic implications to sustainable development. *Journal of Agriculture and Food Research*, 18:101284.
- Faulkner, L. (2003). Beyond the five-user assumption: Benefits of increased sample sizes in usability testing. *Behavior Research Methods, Instruments, & Computers*, 35(3):379–383.
- Food and Agriculture Organization (2023). Knowledge sharing. <https://www.fao.org/knowledge/en/>.
- Global Nutrition Report (2020). Philippines nutrition profile. <https://globalnutritionreport.org/resources/nutrition-profiles/asia/south-eastern-asia/philippines/>.
- Harvard Library (2024). DASH: Digital access to scholarship at Harvard. <https://dash.harvard.edu/>.

- Heitman, K., Hubbard, J., Easter, L., and Kilkus, J. (2024). Looking to the future: Agendas, directions, and resources for nutrition research. *Nutrition in Clinical Practice*.
- Hripcsak, G. and Rothschild, A. S. (2005). Agreement, the f-measure, and reliability in information retrieval. *Journal of the American Medical Informatics Association*, 12(3):296–298.
- IBM (2023). What is knowledge management? <https://www.ibm.com/topics/knowledge-management>.
- International Organization for Standardization (2018). ISO 9241-11:2018 ergonomics of human-system interaction — part 11: Usability: Definitions and concepts. Standard, Geneva, Switzerland.
- Lewis, J. R. (2018). The System Usability Scale: Past, present, and future. *International Journal of Human-Computer Interaction*, 34(7):577–590.
- Lewis, J. R. and Sauro, J. (2018). Item benchmarks for the System Usability Scale. *Journal of Usability Studies*, 13(3):158–167.
- Likert, R. (1932). A technique for the measurement of attitudes. *Archives of Psychology*, 22(140):1–55.
- Maharani, A. et al. (2020). Knowledge management system design using the People, Process, Technology framework. *International Journal of Engineering Research and Technology*.
- McCandless, M., Hatcher, E., and Gospodnetić, O. (2010). *Lucene in Action*. Manning Publications, 2nd edition.
- Nielsen, J. (2000). Why you only need to test with 5 users. <https://www.nngroup.com/articles/why-you-only-need-to-test-with-5-users/>.
- Righettini, M. S. and Bordin, E. (2022). Exploring food security as a multidimensional topic: twenty years of scientific publications and recent developments. *Quality & Quantity*.
- Robertson, S. and Zaragoza, H. (2009). *The Probabilistic Relevance Framework: BM25 and Beyond*. Foundations and Trends in Information Retrieval. Now Publishers.
- Rubin, J. and Chisnell, D. (2008). *Handbook of Usability Testing: How to Plan, Design, and Conduct Effective Tests*. Wiley, 2nd edition.
- Sauro, J. (2011). *A Practical Guide to Measuring Usability: 72 Answers to the Most Common Questions about Quantifying the Usability of Websites and Software*. Measuring Usability LLC.

- Sauro, J. (2012). SEQ: 10 reasons to use the Single Ease Question. <https://measuringu.com/seq10/>.
- Sauro, J. and Dumas, J. S. (2009). Comparison of three one-question, post-task usability questionnaires. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI)*, pages 1599–1608.
- Sauro, J. and Lewis, J. R. (2012). *Quantifying the User Experience: Practical Statistics for User Research*. Morgan Kaufmann.
- Shneiderman, B. (2020). Human-centered artificial intelligence: Reliable, safe and trustworthy. *International Journal of Human-Computer Interaction*, 36(6):495–504.
- Usman, A. et al. (2019). Components of knowledge management systems in modern organizations. *Journal of Information and Knowledge Management*.
- Van Zyl, A. S. et al. (2014). Knowledge management and organizational effectiveness. *Journal of Knowledge Management Practice*.
- Vinutha, M. S. and Padma, M. C. (2020). A survey of AI and ML techniques for search engines. *International Journal of Computer Applications*.